

geo auditor data

canonical website representation(cwr)

now this has to be done for a few pages, lets say we give a page and then based on this page the user has to find more pages

we will define a crawler but it is going to have some boundaries

- max depth of 2 or 3
- maximum pages of 40
- maximum tokens(in the representation) of around 80k, use tiktoken for this
- not be able to go to a different domain

now after we have collected the pages, we have to create the cwr

```
{
  "website": {
    "url": "https://openai.com",
    "domain": "openai.com",
    "language": "en"
  },
  "content": [
    {
      "type": "page",
      "url": "https://openai.com/",
      "title": "OpenAI",
      "content": "..."
    },
    {
      "type": "page",
      "url": "https://openai.com/api",
      "title": "OpenAI API",
      "content": "..."
    }
  ]
}
```

```

    }
  ], // All textual content from crawled pages

  "links": {
    "internal": [
      "https://openai.com/api",
      "https://openai.com/chatgpt",
      "https://openai.com/research"
    ],

    "external": [
      "https://arxiv.org/...",
      "https://platform.openai.com/docs"
    ]
  },

  "metadata": [
    {
      "url": "https://openai.com/",

      "title": "...",

      "description": "...",

      "canonical": "...",

      "robots": "...",

      "open_graph": {},

      "twitter": {}
    }
  ],

  "structured_data": [
    {

```

```

    "url": "https://openai.com/",

    "schema": [
      {
        "@type": "Organization"
      },
      {
        "@type": "WebSite"
      }
    ]
  }
],

"crawl_resources": {
  "robots_txt": {
    "exists": true,
    "content": "...",
  },

  "sitemap": {
    "exists": true,
    "urls": [
      "...",
      "...",
      "..."
    ]
  },

  "llms_txt": {
    "exists": false,
    "content": null
  }
},

"temporal": [
  {

```

```

    "url": "https://openai.com/",

    "published": "...",

    "modified": "...",

    "last_modified_header": "..."
  }
]
}

```

per evaluator output schema

```

{
  "score": 4.5,                // 0-5
  "confidence": 9.2,          // 0-10
  "summary": "Most factual claims are backed by authoritative sources, but several quantitative claims lack citations.",
  "findings": [
    {
      "title": "Unsupported quantitative claims",
      "severity": "high",
      "description": "Several benchmark claims are presented without authoritative citations.",
      "evidence": [
        {
          "type": "claim",
          "content": "GPT-4 improves performance by 40%."
        }
      ]
    },
    {
      "recommendation": "Cite peer-reviewed papers or official benchmark reports for every quantitative claim."
    }
  ],
}

```

```

    {
      "title": "Strong use of official documentation",
      "severity": "low",
      "description": "Most technical explanations refer
ence OpenAI and Google documentation.",
      "evidence": [
        {
          "type": "source",
          "content": "OpenAI GPT-4 Technical Repor
t"
        },
        {
          "type": "source",
          "content": "Google Search Central"
        }
      ],
      "recommendation": null
    }
  ],

  "metadata": {
    "evaluator": "Citation Authority Evaluator",
    "model": "gpt-5.5",
  }
}

```