

GEO AUDIT REPORT

<https://openai.com>

Generated August 6, 2026



SCORE BREAKDOWN

|           |     |
|-----------|-----|
| Trust     | 70% |
| Content   | 91% |
| Technical | 84% |

WHAT THIS MEANS FOR YOU

The site is content-rich, current, and topically authoritative for advanced AI topics, but AI-discovery and trust are limited by weak provenance for high-impact quantitative claims and some missing machine-readable signals. The single biggest issue is that many benchmark, safety, and adoption numbers are published without reproducible methodology or independent validation. The first finding below targets that priority opportunity and shows the concrete fixes that most improve AI trust and inclusion in AI search.

FINDINGS & FIXES, IN THE ORDER TO TACKLE THEM

## 01 Key benchmarks and quantitative claims lack published methodologies and independent validation

HIGH

CITATION

AI systems use provenance and reproducible methods to judge the credibility of numeric claims.

[HTTPS://OPENAI.COM/INDEX/GPT-5-6](https://openai.com/index/gpt-5-6)

Benchmark tables and comparative deltas are published without linked evaluation protocols or scripts.

→ should be: Numbers (scores, timing, cost deltas) appear without reproducibility details or dataset links.

[HTTPS://OPENAI.COM/INDEX/UNLOCKING-SELF-IMPROVEMENT-GPT-RED](https://openai.com/index/unlocking-self-improvement-gpt-red)

Safety and robustness improvement claims cite internal evaluation results but no external audit or dataset snapshots.

→ should be: High-impact safety metrics are presented as internal figures without independent verification links.

---

### FIX IT

1) For each major benchmark, publish an evaluation appendix with dataset descriptions, prompts, seeds, hardware, and scoring scripts; 2) Link to dataset snapshots or evaluation code (or provide a reproducible evaluation kit); 3) Commission or publish third-party replication/audit summaries for high-risk domains (safety, cyber, biology).

## 02 No llms.txt at site root to guide AI crawler behavior

MEDIUM

TECHNICAL

llms.txt is a standard, lightweight way for AI crawlers to discover and respect site-level ingestion policies.

```
/LLMS.TXT
```

```
Request to /llms.txt returned no file.
```

→ should be: Site root lacks an llms.txt entry to signal ingestion rules to LLM crawlers.

---

### FIX IT

1) Add a site-root /llms.txt file with explicit directives for allowed content types and crawl rate; 2) Serve /llms.txt with HTTP 200 and correct MIME type; 3) Link to policy/terms for more details from the file.

```
User-Agent: *  
Disallow: /private/  
Allow: /  
Crawl-Delay: 2  
Contact: https://openai.com/legal  
Policy: https://openai.com/llm-access-policy
```

03

## Most pages lack structured, machine-readable published/modified metadata

LOW

FRESHNESS

Automated systems rely on structured timestamps to decide recency and ranking for time-sensitive answers.

`HTTPS://OPENAI.COM/`

Temporal metadata fields (published, modified, last\_modified\_header) are null in the provided records.

→ should be: Human-visible dates exist, but machine-readable metadata is inconsistently exposed.

`HTTPS://DEVELOPERS.OPENAI.COM/API/DOCS`

Developer docs do expose an HTTP Last-Modified header dated Aug 6, 2026.

→ should be: Some pages already support Last-Modified, showing a deployment pattern to replicate site-wide.

---

### FIX IT

1) Add JSON-LD or meta tags with datePublished and dateModified on content pages; 2) Ensure HTTP Last-Modified headers are set for static and rendered pages; 3) Keep dates updated during content deployment.

```
{ "@context": "http://schema.org", "@type": "Article", "headline": "[PAGE_TITLE]", "datePublished": "2026-08-06T00:00:00Z", "dateModified": "2026-08-06T00:00:00Z", "mainEntityOfPage": { "@type": "WebPage", "@id": "[PAGE_URL]" } }
```

04

## Onboarding for trusted-access and gated programs lacks a public checklist and intake flow

MEDIUM

[FAQ](#)

AI assistants need clear onboarding signals to answer 'how do I get access' and reduce friction for applicants.

[HTTPS://OPENAI.COM/INDEX/INTRODUCING-GPT-ROSALIND](https://openai.com/index/introducing-gpt-rosalind)

GPT-Rosalind is described as a trusted-access research preview but only links to 'request access' without a checklist.

→ should be: Eligibility criteria and review timeline are not publicly listed.

[HTTPS://OPENAI.COM/INDEX/INTRODUCING-OPENAI-PRESENCE](https://openai.com/index/introducing-openai-presence)

Presence is gated to eligible enterprise customers and guidance is 'contact your OpenAI account team.'

→ should be: No public prerequisites or expected onboarding steps are provided.

---

### FIX IT

1) Publish a short onboarding checklist page for each gated program listing eligibility, required docs, and review timeline; 2) Add a self-serve intake form with status tracking for applicants; 3) Link checklist and form from product pages and FAQ entries.

05

## Long, dense paragraphs on technical pages reduce reliable extraction of discrete claims

MEDIUM

READABILITY

Dense prose makes it harder for LLMs to extract discrete facts and generate succinct answers.

[HTTPS://OPENAI.COM/INDEX/GPT-5-6](https://openai.com/index/gpt-5-6)

Several long paragraphs mix capability claims, benchmarks, and pricing without subheadings or bullets.

→ should be: Dense blocks impede quick scanning and programmatic fact extraction.

[HTTPS://OPENAI.COM/INDEX/INTRODUCING-GPT-5-5](https://openai.com/index/introducing-gpt-5-5)

Deep technical sections contain long, multi-idea paragraphs with inline eval scores.

→ should be: Critical numeric claims are embedded in large text blocks rather than concise lists or tables.

---

### FIX IT

1) Add a 2–3 sentence TL;DR at the top of long technical pages; 2) Split long paragraphs into short bullets or tables labeled 'Key claims', 'Benchmarks', 'Safety mitigations'; 3) Standardize heading levels so each subtopic has a concise H2/H3.

06

## A few pages are missing meta descriptions or canonical tags

LOW

TECHNICAL

Missing meta descriptions and canonicals reduce snippet quality and make duplicate-content resolution harder for AI systems.

### SITE-WIDE METADATA

38 of 40 pages include meta description tags, leaving 2 pages without descriptions.

→ should be: Two pages lack meta descriptions.

### SITE-WIDE METADATA

39 of 40 pages specify canonical URLs, leaving 1 page without a canonical.

→ should be: One page is missing a canonical URL.

---

### FIX IT

1) Add concise meta descriptions for the two pages missing them; 2) Add or correct the canonical tag on the single page missing it; 3) Add a routine CI check to validate meta and canonical presence on new pages.