


GEO AUDIT REPORT

<https://sarvam.ai>

Generated August 6, 2026



SCORE BREAKDOWN

Trust	50%
Content	85%
Technical	100%

WHAT THIS MEANS FOR YOU

The site shows strong topical coverage for India-first AI products but AI-facing trust is low because many high-impact claims lack verifiable evidence. The biggest barrier to AI visibility is limited external citations and reproducibility artifacts for benchmarks, usage numbers, and security claims. Fixing external citations and linked evidence will immediately raise machine and human trust. The first finding targets that highest-impact opportunity.

FINDINGS & FIXES, IN THE ORDER TO TACKLE THEM

01 Major quantitative claims lack authoritative, external citations

HIGH

CITATION

LLMs and other AI systems weight external, verifiable citations when judging factual claims.

[HTTPS://WWW.SARVAM.AI/BLOGS/INTRODUCING-SARVAM-CIRCLE](https://www.sarvam.ai/blogs/introducing-sarvam-circle)

"the stack runs at population scale today - more than 2 million voice conversations a day, 35 million pages digitized, roughly 1 crore API calls a day"

→ should be: No external link, audit, or downloadable usage report is provided alongside this large-scale usage claim.

[HTTPS://WWW.SARVAM.AI/](https://www.sarvam.ai/)

"3x increase in customer engagement" and "Millions of policy calls automated"

→ should be: Product highlight metrics are stated with no linked case study or third-party validation.

FIX IT

Add authoritative links for major numeric claims (audit/usage reports, dashboard snapshots, or third-party articles); include a short citation line (source, date, scope) next to every headline metric; surface reproducible appendices or downloadable summaries for auditors.

02 Customer-impact and case metrics are presented without supporting case-study data

HIGH

EVIDENCE

AI systems prefer claims tied to reproducible case data when surfacing factual product outcomes.

[HTTPS://WWW.SARVAM.AI/](https://www.sarvam.ai/)

"50,000+ farmer feedback calls" displayed as a highlight

→ should be: No case-study appendix, baseline, timeframe, or anonymised logs are linked from the claim.

FIX IT

Publish concise case-study appendices per major claim with metric definition, baseline, timeframe, sample size and anonymised aggregate evidence; link each appendix from the homepage and related product pages; add a short methodological note on measurement.

03 Model benchmark results lack reproducibility artifacts and scripts

HIGH

EVIDENCE

Without evaluation scripts, dataset versions, and pinned checkpoints, AI systems cannot verify benchmark claims.

[HTTPS://WWW.SARVAM.AI/BLOGS/SARVAM-30B-105B](https://www.sarvam.ai/blogs/sarvam-30B-105B)

Tabulated benchmark scores (Math500, MMLU, Live Code Bench, etc.) and training token counts

→ should be: No machine-readable evaluation scripts, dataset links, seeds, or model-checkpoint identifiers are linked to reproduce these scores.

FIX IT

Publish a reproducibility repo with evaluation scripts, dataset versions/links, prompts, and pinned model-checkpoint hashes; add a short 'How we evaluated' section linking to that repo on the benchmarks page; include one tiny runnable example notebook.

04 **Certifications and compliance claims are not linked to audit evidence**

HIGH

CITATION

AI systems and enterprises use linked audit artifacts to validate certification claims.

[HTTPS://WWW.SARVAM.AI/TRUST-CENTER](https://www.sarvam.ai/trust-center)

"We're ISO 27001:2022 certified and hold SOC 2 Type II"

→ should be: No scanned certificate, certifier name/ID, or auditor report link is provided on the Trust Center page.

FIX IT

Link scanned certificates or certifier registry entries and include certificate IDs and audit date ranges; provide a short SOC 2 report summary or auditor contact and a download link for certificate artifacts; mark items 'in progress' with expected publication dates.

05 **Training dataset provenance and licensing disclosures are incomplete**

MEDIUM

ENTITY

AI systems and auditors require dataset provenance and license details to validate model claims and reuse constraints.

[HTTPS://WWW.SARVAM.AI/BLOGS/SARVAM-30B-105B](https://www.sarvam.ai/blogs/sarvam-30B-105B)

"trained from scratch on large-scale, high-quality datasets curated in-house"

→ should be: No per-language corpus breakdown, dataset sources, license terms, or filtering/annotation rules are published.

FIX IT

Publish a dataset-provenance summary (per-language corpus sizes, sources, and licenses); add dataset filtering and annotation guidelines; link these items from model-card pages.

06

API and product performance claims lack measurement context and SLA links

MEDIUM

EVIDENCE

AI evaluators and customers need test conditions and SLA references to trust latency and uptime metrics.

[HTTPS://WWW.SARVAM.AI/APIS/TEXT-TO-SPEECH](https://www.sarvam.ai/apis/text-to-speech)

"Sub-250ms first byte with WebSocket streaming"

→ should be: The page does not describe network/test conditions, geographic points, or measurement scripts used to obtain this latency.

[HTTPS://WWW.SARVAM.AI/APIS/SPEECH-TO-TEXT](https://www.sarvam.ai/apis/speech-to-text)

"<250ms Median latency" and ">99.5% Uptime"

→ should be: No SLA document, uptime dashboard link, or measurement window is attached to these availability claims.

FIX IT

Publish measurement methodology and test-harness scripts for latency claims; add a public status/dashboard or SLA summary with measured windows and definitions; link performance dashboards from API pages.

07 Product and docs pages mostly lack publication/last-modified metadata and headers

MEDIUM

FRESHNESS

AI systems use timestamps to assess content currency and to prefer recent authoritative sources.

```
HTTPS://WWW.SARVAM.AI/
```

```
published: null, modified: null, last_modified_header: null
```

→ should be: Homepage and many product/docs pages do not expose visible or machine-readable publish/modified dates.

FIX IT

Add visible datePublished/dateModified meta tags (schema.org) on product and doc pages; return Last-Modified or ETag HTTP headers; maintain a changelog or dated release notes page.

08

Indus product landing pages are empty or gated, blocking public evaluation

MEDIUM

[FAQ](#)

Gated or empty entry points reduce the site's visible surface area for AI systems and users to cite.

`HTTPS://INDUS.SARVAM.AI/`

Page content is empty in the crawl

→ should be: Unauthenticated users hit an empty root or a login redirect that shows no public product summary.

`HTTPS://INDUS.SARVAM.AI/LOGIN?RETURN_TO=%2FTEXT-TO-SPEECH`

"By continuing you agree to our terms of service and privacy policy"

→ should be: Login pages present only an auth notice with no public demo or summary to evaluate product features.

FIX IT

Add a public Indus landing page with a 1–2 sentence summary and demo CTA; surface a short explainer or video for unauthenticated visitors; mark gated sections as 'Private preview' with expected availability.

09

Contact and enterprise support channels are underspecified on contact page

MEDIUM

[FAQ](#)

AI systems and enterprise buyers surface contact info to validate claims and start procurement conversations.

`HTTPS://WWW.SARVAM.AI/CONTACT-US`

Page lists headings like 'Contact Us' and 'Technical support' but no direct contact details

→ should be: No explicit sales email, phone number, enterprise enquiry form, or support-SLA entry is visible in the provided content.

FIX IT

Add explicit enterprise contact channels: sales form, sales/support email and phone, and a clear enterprise onboarding contact; link SLA and onboarding docs on the contact page.

10

A subset of pages are missing canonical URL tags (6 pages identified)

LOW

[TECHNICAL](#)

Canonical tags help crawlers and LLMs deduplicate content and prefer the authoritative URL.

`SITE-WIDE CRAWL METADATA`

34/40 pages specify a canonical URL (i.e., 6 pages missing canonical tags)

→ should be: Ratio indicates multiple pages lack a canonical link element.

FIX IT

Add a `rel="canonical"` link element to each affected page pointing to the preferred URL; ensure canonical points to the canonical homepage or product page as appropriate; verify with a post-deploy crawl.

```
<link rel="canonical"
href="https://www.sarvam.ai/REPLACE_WITH_PAGE_PATH" />
```


11 Homepage and marketing hero blocks rely on short slogans without connective explanation LOW READABILITY

AI extractors and users need one-line context to map slogans to concrete capabilities.

`HTTPS://WWW.SARVAM.AI/`

`"India's Sovereign AI Platform / Built on sovereign compute.
Powered by frontier-class models."`

→ should be: Hero contains short slogans without a 1–2 sentence description of what the product actually does and who it's for.

FIX IT

Add a 1–2 sentence explanatory line under hero and key marketing blocks that states product purpose and target audience; ensure that explanatory lines are present in machine-readable HTML near the heading; keep each explanatory sentence under 25 words.

12 Model release text claims weights are available but does not link direct artifact pages or checksums LOW CITATION

AI systems and researchers need direct artifact links and checksums to validate open-weight claims.

`HTTPS://WWW.SARVAM.AI/BLOGS/SARVAM-30B-105B`

`"Weights can be downloaded from AI Kosh (30B,105B) and Hugging
Face (30B,105B)"`

→ should be: Blog text names distribution channels but does not include clickable links, model-card pages, or checksum identifiers in the excerpt.

FIX IT

Add direct links to model artifact pages (AI Kosh, Hugging Face) and model-card pages; publish checksums or release tags and link to a public release notes page; include brief usage examples and inference costs per model.

GEO Auditor · report generated automatically · verify before publishing